

Done

```
import numpy as np
import pandas as pd
import re
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

import nltk
nltk.download('stopwords')
print(stopwords.words('english'))

['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'you
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
td=pd.read_csv('/content/training.1600000.processed.noemoticon.csv', encoding="ISO-8859-1")
```

```
td.shape

(1599999, 6)
```

```
td.head()
```

			Mon Apr 06 22:19:45 PDT 2009	NO_QUERY	_TheSpecialOne_	@switchfoot http://twitpic.com/2y1z - Awww, that's a bumme You shoul Carr of Third Day to it.
0	0	1467810369	Mon Apr 06 22:19:49 PDT 2009	NO_QUERY	scotthamilton	is upset that he can't update h Facebook by
1	0	1467810917	Mon Apr 06 22:19:53 PDT 2009	NO_QUERY	mattycus	@Kenichan I dived many tim for the ball. Mar
			Mon Apr			

```
cn=['target', 'id', 'date', 'flag', 'user', 'text']
td=pd.read_csv('/content/training.1600000.processed.noemoticon.csv',names=cn, encoding="ISO-8859-1")
```

```
td.shape

(1600000, 6)
```

```
td.head()
```

		target	id	date	flag	user	text
0	0		1467810369	Mon Apr 06 22:19:45 PDT 2009	NO_QUERY	_TheSpecialOne_	@switchfoot http://twitpic.com/2y1zl - Awww, t...
1	0		1467810672	Mon Apr 06 22:19:49 PDT 2009	NO_QUERY	scotthamilton	is upset that he can't update his Facebook by ...

```
td.isnull().sum()
```

```
target    0
id        0
date      0
flag      0
user      0
text      0
dtype: int64
```

```
td['target'].value_counts()
```

```
0    800000
4    800000
Name: target, dtype: int64
```

```
td.replace({'target': {4:1}},inplace=True)
td['target'].value_counts()
```

```
0    800000
1    800000
Name: target, dtype: int64
```

```
ps=PorterStemmer()
```

```
def stemming(sc):
    sc = re.sub('[^a-zA-Z]', ' ', sc)
    sc = sc.lower()
    sc = sc.split()
    sc = [ps.stem(word) for word in sc if word not in stopwords.words('english')]
    sc = ' '.join(sc)
    return sc
```

```
#td['stemmed_content'] = td['text'].apply(stemming)
```

```
import pickle
```

```
td['sc']=pickle.load(open('/content/stemmed_model.sav', 'rb'))
```

```
print(td['sc'])
```

```
0    switchfoot http twitpic com zl awww bummer sho...
1    upset updat facebook text might cri result sch...
2    kenichan dive mani time ball manag save rest g...
3    whole bodi feel itchi like fire
4    nationwideclass behav mad see
...
1599995    woke school best feel ever
1599996    thewdb com cool hear old walt interview http b...
1599997    readi mojo makeov ask detail
1599998    happi th birthday boo alll time tupac amaru sh...
1599999    happi charitytuesday thenspcc sparksschar speak...
Name: sc, Length: 1600000, dtype: object
```

```
print(td['target'])
```

```
0    0
1    0
2    0
3    0
4    0
...
1599995    1
1599996    1
1599997    1
1599998    1
1599999    1
Name: target, Length: 1600000, dtype: int64
```

```
x=td['sc'].values
```

```
y=td['target'].values
```

```
print(x)
```

```

['switchfoot http twitpic com zl awww bummer shoulda got david carr third day'
'upset updat facebook text might cri result school today also blah'
'kenichan dive mani time ball manag save rest go bound' ...
'readi mojo makeov ask detail'
'happi th birthday boo alll time tupac amaru shakur'
'happi charitytuesday thenspcc sparksschar speakinguph h']

print(y)

[0 0 0 ... 1 1 1]

x_train,x_test,y_train,y_test=train_test_split(x, y, test_size=0.3, stratify=y, random_state=2)

print(x.shape, x_train.shape, x_test.shape)

(1600000,) (1120000,) (480000,)

v=TfidfVectorizer()

x_train=v.fit_transform(x_train)

x_test=v.transform(x_test)

print(x_train)

(0, 374778) 0.31949818170660904
(0, 328004) 0.5314491135690684
(0, 43969) 0.4748653389926461
(0, 306319) 0.4371164314704369
(0, 121348) 0.44599875194306554
(1, 360908) 0.20232270128356875
(1, 45367) 0.24454930105105102
(1, 32222) 0.2885646935506797
(1, 400139) 0.5008486447803013
(1, 36162) 0.21760975021723114
(1, 376966) 0.16230856596272633
(1, 224362) 0.24716194206021685
(1, 2941) 0.3329517305729671
(1, 28528) 0.1699742942038265
(1, 377777) 0.5390797727604116
(2, 209149) 0.36506885244060977
(2, 217307) 0.35862123416780045
(2, 331146) 0.45772430309259293
(2, 192775) 0.7270515839609045
(3, 54144) 0.3405985100366321
(3, 381278) 0.3044716859774938
(3, 114004) 0.3165169852577187
(3, 344001) 0.3565976148233519
(3, 366377) 0.4018889828996008
(3, 139076) 0.25759117714273316
:
(1119996, 327788) 0.2626904633863284
(1119996, 133766) 0.14165568151908714
(1119996, 400064) 0.17675763929108915
(1119996, 366962) 0.1645582384006955
(1119996, 121102) 0.2083143336548096
(1119996, 166244) 0.2582309496394417
(1119997, 364520) 0.7210223776952688
(1119997, 176159) 0.6575712947692484
(1119997, 366962) 0.21846446658016408
(1119998, 384964) 0.23515032977930925
(1119998, 26391) 0.22101418705941295
(1119998, 255402) 0.23115328947076547
(1119998, 313788) 0.49883995385230084
(1119998, 405750) 0.597655663753657
(1119998, 252433) 0.34998949156158465
(1119998, 56066) 0.20742679758052754
(1119998, 306613) 0.14119239340870796
(1119998, 155263) 0.14334431819822027
(1119998, 36162) 0.17432788141763644
(1119999, 88171) 0.5404842471107606
(1119999, 124023) 0.61479006915859
(1119999, 397207) 0.36018431710547344
(1119999, 28754) 0.24841651195142356
(1119999, 399014) 0.31666662934784157
(1119999, 221882) 0.19541926631660167

```

```

model = LogisticRegression(max_iter=1000)

model.fit(x_train, y_train)

y_train_pred = model.predict(x_train)

y_test_pred = model.predict(x_test)

training_accuracy = accuracy_score(y_train, y_train_pred)

testing_accuracy = accuracy_score(y_test, y_test_pred)

print(f'Training Accuracy: {training_accuracy}')
print(f'Testing Accuracy: {testing_accuracy}')

    Training Accuracy: 0.8106169642857143
    Testing Accuracy: 0.7776708333333333

from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(y_test, y_test_pred)
recall = recall_score(y_test, y_test_pred)
f1 = f1_score(y_test, y_test_pred)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")

    Precision: 0.7671
    Recall: 0.7974
    F1 Score: 0.7820

def predict_sentiment(text):

    text = stemming(text)

    text_vector = v.transform([text])

    sentiment = model.predict(text_vector)

    return 'Positive' if sentiment == 1 else 'Negative'

new_sentence = "Movie is Good to watch";

print(f"The sentiment of the sentence is: {predict_sentiment(new_sentence)}")

    The sentiment of the sentence is: Positive

from sklearn.naive_bayes import MultinomialNB
m2 = MultinomialNB()
m2.fit(x_train, y_train)
x_train_pred = m2.predict(x_train)
trd_acc = accuracy_score(y_train, x_train_pred)
print('Accuracy on Training Data:', trd_acc)
x_test_pred = m2.predict(x_test)
tsd_acc = accuracy_score(y_test, x_test_pred)
print('Accuracy on Test Data:', tsd_acc)

    Accuracy on Training Data: 0.8216964285714285
    Accuracy on Test Data: 0.7554416666666667

from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, x_test_pred)
recall = recall_score(y_test, x_test_pred)
f1 = f1_score(y_test, x_test_pred)

print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'F1 Score: {f1}')

    Precision: 0.7709082713943561
    Recall: 0.7268958333333333

```

F1 Score: 0.7482554074861998

```
from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(random_state=2)
sgd_clf.fit(x_train, y_train)
x_train_pred = sgd_clf.predict(x_train)
trd_acc = accuracy_score(y_train, x_train_pred)
print('Accuracy on Training Data:', trd_acc)
```

```
x_test_pred = sgd_clf.predict(x_test)
tsd_acc = accuracy_score(y_test, x_test_pred)
print('Accuracy on Test Data:', tsd_acc)
```

```
precision = precision_score(y_test, x_test_pred)
recall = recall_score(y_test, x_test_pred)
f1 = f1_score(y_test, x_test_pred)
```

Accuracy on Training Data: 0.75661875  
Accuracy on Test Data: 0.7561770833333333

```
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'F1 Score: {f1}')
```

Precision: 0.7271165716692094  
Recall: 0.8201541666666666  
F1 Score: 0.7708381876959286

```
from sklearn.svm import SVC
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score
svc_clf = SVC(random_state=2)
svc_clf.fit(x_train, y_train)
```